

Deep Learning 101

Convolutional Neural Network (CNN)

Schedule

week	Date	Topic
9	10.27	Environment setup, python, Jupyter, PyCharm, TensorFlow, & regression
10	11.03	Training and testing
11	11.11	CNN
12	11.18	RNN
13	11.24	Autoencoder & GAN

Today's Class

- Recap
- What is CNN?
 - What is convolution?
 - Stride
 - Padding
 - Max pooling
 - Dropout
- Confusion metrics
- Lab time

Recap – neural network

- Neural network as a function
 - $y = f(x)$
- Perceptron
 - $Y = WX + b$
 - Two inputs: x_1, x_2
 - One output: y
 - Linear regression
- XOR problem
 - Linear regression can't solve the XOR problem
 - Require multivariate regression

Recap - Training

- What it take to train a neural network:
 - Hypothesis: $H = WX + b$
 - Activation function: Sigmoid, tanh, ReLU, LeakyReLU, Softmax, etc.
 - Cost function: MSE, cross entropy, etc.
 - Gradient descent: backpropagation
- Training a neural network is basically the problem of minimizing the cost function: minimize $\text{cost}(W, b)$
- Gradient descent is the most popular optimizer.
- Training a neural network is NOT easy!
 - Finding hyperparameters, random initial weights, local minima, vanishing/exploding gradients, overfitting/underfitting, etc.

How does convolutions work?

3x3 portion of an image		filter																			
<table><tr><td>0.6</td><td>0.2</td><td>0.6</td></tr><tr><td>0.1</td><td>-0.2</td><td>-0.3</td></tr><tr><td>-0.5</td><td>-0.1</td><td>-0.3</td></tr></table>	0.6	0.2	0.6	0.1	-0.2	-0.3	-0.5	-0.1	-0.3	*	<table><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>-1</td><td>-1</td><td>-1</td></tr></table>	1	1	1	0	0	0	-1	-1	-1	= 2.3
0.6	0.2	0.6																			
0.1	-0.2	-0.3																			
-0.5	-0.1	-0.3																			
1	1	1																			
0	0	0																			
-1	-1	-1																			
<table><tr><td>-0.6</td><td>-0.2</td><td>-0.6</td></tr><tr><td>-0.1</td><td>0.2</td><td>0.3</td></tr><tr><td>0.5</td><td>0.1</td><td>0.3</td></tr></table>	-0.6	-0.2	-0.6	-0.1	0.2	0.3	0.5	0.1	0.3	*	<table><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>-1</td><td>-1</td><td>-1</td></tr></table>	1	1	1	0	0	0	-1	-1	-1	= -2.3
-0.6	-0.2	-0.6																			
-0.1	0.2	0.3																			
0.5	0.1	0.3																			
1	1	1																			
0	0	0																			
-1	-1	-1																			


* Generative Deep Learning (David Foster)

What is involved in convolution

- Stride: step size
- Padding: putting zeroes around the outer edge of the input data.
“same” means the output size will be the same as the input size when stride = 1.
- Kernel: the filter that extracts features
- Max pooling: means pooling the maximum value from the filtered feature map. The result is a down-sampling image (reduced dimensionality)
- Batch normalization: normalizing
- Dropout: regularization technique to prevent overfitting


Convolution in action

- stride = 1, padding = “same”




- What will happen if there is no padding?

Max pooling




Max pooling with 2x2
filters with stride 2




What normalization means

- Zero-centered normalized:



Batch normalization in Deep Learning (intuition)




Batch normalization

- To standardize the input: to mitigate the "covariate shift" problems

Batch normalization

Ensure the output statistics of a layer are fixed.



Batch normalization in action

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

<https://www.jeremyjordan.me/batch-normalization/>

- normalize the activations of the previous layer for each batch
- apply a transformation that maintains the mean activation close to 0 and the activation standard deviation close to 1.

Visualization Tools

- [Convolution Visualizer](#)
- [CNN Explainer](#)

Confusion matrix


		Actual	
		Positive	Negative
Predicted	Positive	True Positive	False Positive
	Negative	False Negative	True Negative

- Are you confused?

Actual example: COVID-19

		Actual	
		Yes	No
Predicted	Yes	True Positive	False Positive
	No	False Negative	True Negative

- Which is worse? False Negative or False Positive?
- What if the problem is different? e.g., spam filter.



How many selected items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

Metrics: accuracy, precision, and recall

$$\text{Accuracy} = \frac{\text{True Positives} + \text{True Negatives}}{\text{All Samples}}$$

$$\textit{Precision} = \frac{\textit{TruePositive}}{\textit{TruePositive} + \textit{FalsePositive}}$$

$$\textit{Recall} = \frac{\textit{TruePositive}}{\textit{TruePositive} + \textit{FalseNegative}}$$

Metrics

- Problems with accuracy: imbalanced classification problems
- Precision: what proportion of positive predictions was actually correct?
 - $(TP) / (TP+FP)$
- Recall: what proportion of actual positives was predicted correctly?
 - $(TP) / (TP+FN)$
-

Lab time

- To clone: from your terminal
 - >git clone <https://github.com/changsin/DeepLearning-101.git>
- Or use google colab to point to the git hub repository
- Git is an open source version control system
 - Github is a host service using git.